

Before you turn in the homework, make sure everything runs as expected. To do so, select **Kernel**→**Restart & Run All** in the toolbar above. Remember to submit both on **DataHub** and **Gradescope**.

Please fill in your name and include a list of your collaborators below.

In [1]:

```
1 NAME = "Junsheng Pei"
2 COLLABORATORS = ""
```

## Project 2: NYC Taxi Rides

### Extras

Put all of your extra work in here. Feel free to save figures to use when completing Part 4.

In [2]:

```
1 import os
2 import pandas as pd
3 import numpy as np
4 from pathlib import Path
5 from sqlalchemy import create_engine
6 from utils import timeit
7 import matplotlib.pyplot as plt
8 import seaborn as sns
```

In [3]:

```
1 !ls -lh /srv/db/taxi_2016_student_small.sqlite

-rw-r--r-- 1 root root 2.1G Nov  7 04:44 /srv/db/taxi_2016_student_small.sqlite
```

In [4]:

```
1 DB_URI = "sqlite:///srv/db/taxi_2016_student_small.sqlite"
2 TABLE_NAME = "taxi"
3
4 sql_engine = create_engine(DB_URI)
5 with timeit():
6     print(f"Table {TABLE_NAME} has {sql_engine.execute(f'SELECT COUNT(*) FROM {TABLE_NAME}').first()[0]} rows!")
```

Table taxi has 15000000 rows!  
1.15 s elapsed

### Data Selection:

we chose data from April instead of January, and then we export the data

In [5]:

```
1 query = f"""
2     SELECT *
3     FROM (
4     SELECT *
5     FROM (
6     SELECT *
7     FROM (
8     SELECT *
9     FROM taxi
10    WHERE tpep_pickup_datetime
11           BETWEEN '2016-04-01' AND '2016-04-30'
12           AND record_id % 100 == 0
13    ORDER BY tpep_pickup_datetime
14    )
15    WHERE (julianday(tpep_dropoff_datetime) - julianday(tpep_pickup_datetime)) < 0.5
16    )
17    )
18    WHERE passenger_count > 0
19    """
20
21 with timeit(): # this query should take less than a second
22     cleaned_df = pd.read_sql(query, sql_engine)
23 cleaned_df['tpep_pickup_datetime'] = pd.to_datetime(cleaned_df['tpep_pickup_datetime'])
24 cleaned_df['tpep_dropoff_datetime'] = pd.to_datetime(cleaned_df['tpep_dropoff_datetime'])
25 cleaned_df['duration'] = cleaned_df["tpep_dropoff_datetime"]-cleaned_df["tpep_pickup_datetime"]
26 cleaned_df['duration'] = cleaned_df['duration'].dt.total_seconds()
```

2.62 s elapsed

In [6]:

```
1 from sklearn.model_selection import train_test_split
2 train_df, val_df = train_test_split(cleaned_df, test_size=0.2, random_state=42)
3 data_file = Path(".", "cleaned_data_2016.hdf") # Path of hdf file
4 train_df.to_hdf(data_file, "train") # Train data of hdf file
5 val_df.to_hdf(data_file, "val") # Val data of hdf file
```

### Better features

We test the the following features: 'tip\_amount', 'trip\_distance', 'ifdaytime', 'ifweekday', 'total\_amount', 'fare\_amount.


And we found those feature can help us predict duration

In [7]:


```
1 cleaned_df = cleaned_df[cleaned_df['duration'] < 10000]
```

In [8]:


```
1 cleaned_df_feature = cleaned_df[cleaned_df['tip_amount'] < 40]
2 sns.scatterplot('tip_amount', 'duration', data = cleaned_df_feature)
3 plt.title('tips')
4 plt.show()
```




```
In [9]: 1 cleaned_df_feature = cleaned_df[cleaned_df['tip_amount'] > 0]
2 sns.scatterplot('total_amount', 'duration', data = cleaned_df_feature)
3 plt.title('total_amount')
4 plt.show()
```



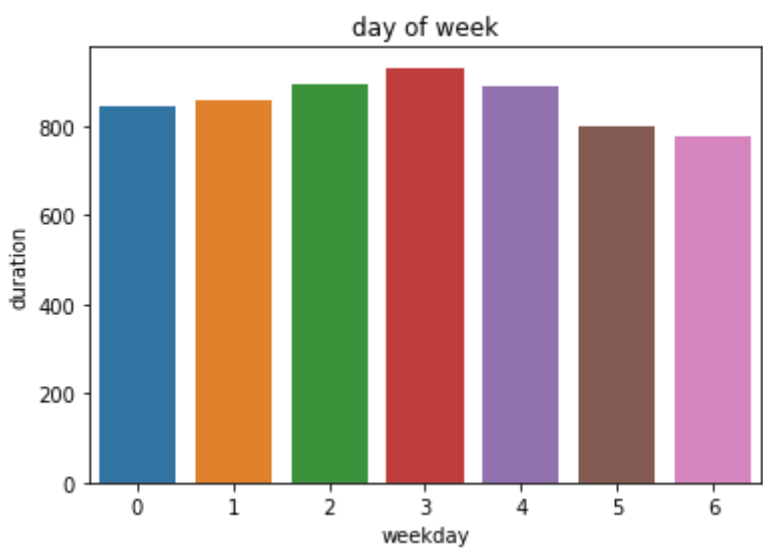
```
In [10]: 1 cleaned_df_feature = cleaned_df[cleaned_df['fare_amount'] > 0]
2 sns.scatterplot('fare_amount', 'duration', data = cleaned_df_feature)
3 plt.title('fare_amount')
4 plt.show()
```




```
In [11]: 1 cleaned_df_feature = cleaned_df[cleaned_df['trip_distance'] > 0]
2 sns.scatterplot('trip_distance', 'duration', data = cleaned_df_feature)
3 plt.title('trip_distance')
4 plt.show()
```



```
In [12]: 1 cleaned_df['weekday'] = cleaned_df['tpep_pickup_datetime'].dt.dayofweek
2 groupByweekday = cleaned_df.groupby('weekday')['duration'].mean()
3 sns.barplot(x =groupByweekday.index, y = groupByweekday)
4 plt.title('day of week')
5 plt.show()
```



```
In [13]: 1 cleaned_df['hour'] = cleaned_df['tpep_pickup_datetime'].dt.hour
2 groupByhour = cleaned_df.groupby('hour')['duration'].mean()
3 sns.barplot(x =groupByhour.index, y = groupByhour)
4 plt.title('day of hour')
5 plt.show()
```



## Feature Selection and Data Cleaning

```
In [14]: 1 data_file_fm = Path("./", "cleaned_data_2016.hdf")
2 train_df_fm = pd.read_hdf(data_file_fm, "train")
3 val_df_fm = pd.read_hdf(data_file_fm, "val")
```

```
In [15]: 1 # Copied from part 2
2 def haversine(lat1, lng1, lat2, lng2):
3     """
4     Compute haversine distance
5     """
6     lat1, lng1, lat2, lng2 = map(np.radians, (lat1, lng1, lat2, lng2))
7     average_earth_radius = 6371
8     lat = lat2 - lat1
9     lng = lng2 - lng1
10    d = np.sin(lat * 0.5) ** 2 + np.cos(lat1) * np.cos(lat2) * np.sin(lng * 0.5) ** 2
11    h = 2 * average_earth_radius * np.arcsin(np.sqrt(d))
12    return h
13
14 # Copied from part 2
15 def manhattan_distance(lat1, lng1, lat2, lng2):
16     """
17     Compute Manhattan distance
18     """
19     a = haversine(lat1, lng1, lat1, lng2)
20     b = haversine(lat1, lng1, lat2, lng1)
21     return a + b
22
23 # Copied from part 2
24 def bearing(lat1, lng1, lat2, lng2):
25     """
26     Compute the bearing, or angle, from (lat1, lng1) to (lat2, lng2).
27     A bearing of 0 refers to a NORTH orientation.
28     """
29     lng_delta_rad = np.radians(lng2 - lng1)
30     lat1, lng1, lat2, lng2 = map(np.radians, (lat1, lng1, lat2, lng2))
31     y = np.sin(lng_delta_rad) * np.cos(lat2)
32     x = np.cos(lat1) * np.sin(lat2) - np.sin(lat1) * np.cos(lat2) * np.cos(lng_delta_rad)
33     return np.degrees(np.arctan2(y, x))
34
35 # Copied from part 2
36 def add_time_columns(df):
37     """
38     Add temporal features to df
39     """
40     df.is_copy = False # propogate write to original dataframe
41     df.loc[:, 'month'] = df['tpep_pickup_datetime'].dt.month
42     df.loc[:, 'week_of_year'] = df['tpep_pickup_datetime'].dt.weekofyear
43     df.loc[:, 'day_of_month'] = df['tpep_pickup_datetime'].dt.day
44     df.loc[:, 'day_of_week'] = df['tpep_pickup_datetime'].dt.dayofweek
45     df.loc[:, 'hour'] = df['tpep_pickup_datetime'].dt.hour
46     df.loc[:, 'week_hour'] = df['tpep_pickup_datetime'].dt.weekday * 24 + df['hour']
47     return df
48
49 # Copied from part 2
50 def add_distance_columns(df):
51     """
52     Add distance features to df
53     """
54     df.is_copy = False # propogate write to original dataframe
55     df.loc[:, 'manhattan'] = manhattan_distance(lat1=df['pickup_latitude'],
56                                                  lng1=df['pickup_longitude'],
57                                                  lat2=df['dropoff_latitude'],
58                                                  lng2=df['dropoff_longitude'])
59
60     df.loc[:, 'bearing'] = bearing(lat1=df['pickup_latitude'],
61                                   lng1=df['pickup_longitude'],
62                                   lat2=df['dropoff_latitude'],
63                                   lng2=df['dropoff_longitude'])
64     df.loc[:, 'haversine'] = haversine(lat1=df['pickup_latitude'],
65                                       lng1=df['pickup_longitude'],
66                                       lat2=df['dropoff_latitude'],
67                                       lng2=df['dropoff_longitude'])
68     return df
69
70 def select_columns(data, *columns):
71     return data.loc[:, columns]
72
73 def mae(actual, predicted):
74     """
75     Calculates MAE from actual and predicted values
76     Input:
77         actual (1D array-like): vector of actual values
78         predicted (1D array-like): vector of predicted/fitted values
79     Output:
80         a float, the MAE
81     """
82
83     mae = np.mean(np.abs(actual - predicted))
84     return mae
```

```
In [16]: 1 def add_ifdaytime(data):
2     data['ifdaytime'] = (data['hour'] >= 8) & (data['hour'] <= 18)
3     return data
4
5 def add_ifweekday(data):
6     data['ifweekday'] = data['day_of_week'] > 4
7     return data
8
9 def drop_outlier(data, col, _filter):
10    return data.loc[data[col][lambda x: _filter(x)].index]
11
12 def replace_outlier(data, col, _filter):
13    mean = data[col][lambda x: _filter(x)].mean()
14    data[col] = data[col].apply(lambda x : x if _filter(x) else mean)
15    return data
16
```

```
In [17]: 1 def process_data_fm(data, test=False):
2         # Put your final pipeline here
3
4         # data cleaning
5         if(test):
6             clean_data = replace_outlier
7         else:
8             clean_data = drop_outlier
9             data = clean_data(data, 'duration', lambda x : (x < 8000) & x > 0)
10
11         filter_latitude = lambda x : (x >= 40.63) & (x <= 40.85)
12
13         data = clean_data(data, 'pickup_latitude', filter_latitude )
14         data = clean_data(data, 'dropoff_latitude', filter_latitude )
15
16         filter_longitude = lambda x : (x >= -74.03) & (x <= -73.75)
17
18         data = clean_data(data, 'pickup_longitude', filter_longitude)
19         data = clean_data(data, 'dropoff_longitude', filter_longitude )
20
21         data = clean_data(data, 'total_amount', lambda x: (x>0) & (x <= 90))
22         data = clean_data(data, 'fare_amount', lambda x: (x>0) & (x <= 80))
23         data = clean_data(data, 'tip_amount', lambda x :(x>0) & (x <= 20) )
24
25         data = clean_data(data, 'trip_distance', lambda x : x < 50)
26
27         X = (
28             data
29             # Transform data
30             .pipe(add_time_columns)
31             .pipe(add_distance_columns)
32             .pipe(add_ifdaytime)
33             .pipe(add_ifweekday)
34             .pipe(select_columns,
35                 'pickup_longitude',
36                 'pickup_latitude',
37                 'dropoff_longitude',
38                 'dropoff_latitude',
39                 'manhattan',
40                 'tip_amount',
41                 'haversine',
42                 'trip_distance',
43                 'ifdaytime',
44                 'ifweekday',
45                 'total_amount',
46                 'fare_amount',
47             )
48         )
49         if test:
50             y = None
51         else:
52             y = data['duration']
53
54         return X, y
55
56 # YOUR CODE HERE
57 #raise NotImplementedError()
```

## Parameter Section for Ridge and Lasso and Model Selection

```
In [18]: 1 import sklearn.linear_model as lm
```

```
In [19]: 1 # Parameter Section for Ridge
2 _lambdas = [0.01,0.1,0.2,0.5,1,2,3,5,10]
3 X_train_fm, y_train_fm = process_data_fm(train_df_fm)
4 X_val_fm, y_val_fm = process_data_fm(val_df_fm)
5
6 #final_model = lm.LinearRegression(fit_intercept=True)
7 for _lambda in _lambdas:
8     final_model = lm.Ridge(alpha = _lambda, fit_intercept=True)
9
10    # Define your final model here, feel free to try other forms of regression
11    final_model.fit(X_train_fm, y_train_fm)
12
13    y_val_pred_fm = final_model.predict(X_val_fm)
14
15    print("for lambda " + str(_lambda) + ", validation accuracy: " + str(mae(y_val_pred_fm,y_val_fm)))
```

```
/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generic.py:4388: FutureWarning: Attribute 'is_copy' is deprecated and will be removed in a future ver
sion.
    object.__getattr__(self, name)
/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generic.py:4389: FutureWarning: Attribute 'is_copy' is deprecated and will be removed in a future ver
sion.
    return object.__setattr__(self, name, value)
```

```
for lambda 0.01, validation accuracy: 122.773109937
for lambda 0.1, validation accuracy: 122.715747715
for lambda 0.2, validation accuracy: 122.655868278
for lambda 0.5, validation accuracy: 122.502909548
for lambda 1, validation accuracy: 122.325242267
for lambda 2, validation accuracy: 122.109814835
for lambda 3, validation accuracy: 122.027044081
for lambda 5, validation accuracy: 122.060878093
for lambda 10, validation accuracy: 122.361729584
```



```
In [20]: 1 # Parameter Section for Ridge
2 _lambdas = [0.01,0.1,0.2,0.5,1,2,3,5,10]
3 X_train_fm, y_train_fm = process_data_fm(train_df_fm)
4 X_val_fm, y_val_fm = process_data_fm(val_df_fm)
5
6 #final_model = lm.LinearRegression(fit_intercept=True)
7 for _lambda in _lambdas:
8     final_model = lm.Lasso(alpha = _lambda, fit_intercept=True)
9
10     # Define your final model here, feel free to try other forms of regression
11     final_model.fit(X_train_fm, y_train_fm)
12
13     y_val_pred_fm = final_model.predict(X_val_fm)
14
15     print("for lambda " + str(_lambda) + ", validation accuracy: " + str(mae(y_val_pred_fm,y_val_fm)))
```

/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generic.py:4388: FutureWarning: Attribute 'is\_copy' is deprecated and will be removed in a future version.

object.\_\_getattr\_\_(self, name)
/srv/conda/envs/data100/lib/python3.6/site-packages/pandas/core/generic.py:4389: FutureWarning: Attribute 'is\_copy' is deprecated and will be removed in a future version.

return object.\_\_setattr\_\_(self, name, value)
/srv/conda/envs/data100/lib/python3.6/site-packages/sklearn/linear\_model/coordinate\_descent.py:491: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Fitting data with very small alpha may cause precision problems.
ConvergenceWarning)

for lambda 0.01, validation accuracy: 122.720904993
for lambda 0.1, validation accuracy: 122.359093687
for lambda 0.2, validation accuracy: 122.2259669
for lambda 0.5, validation accuracy: 122.37614613
for lambda 1, validation accuracy: 124.429678243
for lambda 2, validation accuracy: 124.532659775
for lambda 3, validation accuracy: 124.788272559
for lambda 5, validation accuracy: 125.625109785
for lambda 10, validation accuracy: 128.680014266

So we find using Ridge regression with lambda = 3 is the best

## Submission

You're almost done!

Before submitting this assignment, ensure that you have:

1. Restarted the Kernel (in the menubar, select Kernel→ Restart & Run All)
2. Validated the notebook by clicking the "Validate" button.

Then,

1. **Submit** the assignment via the Assignments tab in **Datahub**
2. **Upload and tag** the manually reviewed portions of the assignment on **Gradescope**